

SFS Lab: Writing a Multithreaded File System

15-213/15-513: Introduction to Computer Systems — self-study edition

Last updated: September 4, 2026

This self-study edition independently adapts the Shark File System (SFS) lab from CMU 15-213/15-513 (Summer 2024). Original materials and lab design remain the property of Carnegie Mellon University and course staff. This is not an official release; keep solutions private.

1 Introduction

A file system is a data structure that the operating system uses to store, maintain, and refer to meaningful data on the disk, and conversely to derive usable information from the raw data on the disk. File systems allow programs to interact with data as abstract files, leaving the cumbersome work of putting that data onto the physical disk to lower-level calls. File systems are also simple but powerful tools for organization (with folders, linking, etc.) and security (with permissions).

In this lab, we provide you with a very basic file system, which does not implement folders, linking, or permissions. In the Shark File System, as distributed in this lab, you may only create and delete files, and read or write to them.

The lab has two parts. The first part deals with file system capabilities: you will implement a few functions that introduce new abilities of the file system, allowing users to see and change the position of file descriptors, and to rename files with a single call.

The second part asks you to improve the file system's safety and performance. Currently, there is nothing to prevent two simultaneous calls to the file system from overwriting or otherwise conflicting with each other. You will prevent conflicts like these from happening, while still allowing concurrent and fast access to the file system for separate users.

2 Logistics

This edition of the lab is self-contained and self-graded: there is no Autolab, no handin directory, and no course server. The test driver described in Section 6 plays the role of the autograder, and the score it prints is local feedback, not an official CMU grade.

The scored student work remains confined to `sfs-disk.c`; do not modify `sfs-api.h`. Optional routes may also require synchronized changes to `sfs-disk.h` and `sfs-fsck.c`.

The lab is intended for a 64-bit GNU/Linux environment (a native installation, WSL2, or a Linux container) with `gcc`, `make`, and POSIX threads.

3 Handout Instructions

Start by extracting the handout tarball in a directory where you plan to do your work:

```
linux> tar xf sfslab-handout.tar
linux> cd sfslab
linux> make
```

The important files are:

`sfs-disk.c`

The SFS implementation, and the only file you need to edit. Three functions in it are unimplemented, and the whole file is thread-unsafe until you make it otherwise.

`sfs-api.h`

The API specification. The doc comment above each function tells you exactly what that function must do. Read it carefully before writing any code.

`sfs-disk.h`

The on-disk data structures and constants. Worth reading: it defines the high-level design of the file system.

`sfs-support.c`

Low-level disk image and mmap support routines. Provided; do not modify.

`test-sfs.c`

The test driver and autograder (Section 6).

`sfs-fsck.c`

A file system consistency checker (Section 7.2).

`sfs-baseline-ref.c`

The handout implementation wrapped in a single global mutex. Used by `make baseline` to calibrate the performance score for your machine. Do not modify.

Typing `make` produces three executables: `test-sfs` (the driver), `sfs-fsck` (the checker), and `test-sfs-baseline` (the calibration reference). Before you implement anything, run `./test-sfs` once: traces A00, A01, B00, and all of Category C should already pass, and the traces that exercise `sfs_getpos`, `sfs_seek`, and `sfs_rename` should fail. That is the expected starting state, and your job is to make the rest of the scoreboard green.

4 The SFS File System

4.1 The on-disk format

An SFS disk image is an ordinary file, mapped into memory with `mmap(2)` and divided into 512-byte blocks. Block 0 is the *superblock*; it identifies the image (a magic number), records the total number of blocks, and embeds the root directory: an array of 15 fixed-size directory entries, each holding a file name (at most 23 characters plus a terminating NUL), the file's size in bytes, and the block number of the file's first block. A directory entry whose first-block field is zero is unused.

Every other block begins with a 12-byte header containing a type tag and the block numbers of the previous and next blocks in whatever chain the block belongs to. The remaining 500 bytes of each block hold file data. Blocks form doubly linked lists: every file is a chain of blocks, and all unallocated blocks sit on a free list rooted in the superblock. Block number 0 acts as the NULL pointer, so the superblock can never appear in the middle of a chain. This is the same flavor of design as the classic FAT file system, except that the links live inside the blocks themselves.

Every file owns at least one block; a nonzero first-block field marks its directory entry live. Non-empty files use one block per 500 bytes. The size field and the last block's zero next pointer both mark the end of a file.

4.2 The provided API

The public entry points, declared and specified in `sfs-api.h`, are:

```
sfs_format(diskName, diskSize)
    Create and format a fresh disk image, then make it the active image.

sfs_mount(diskName) / sfs_unmount()
    Load an existing image / write out and deactivate the current one. Unmounting fails with -EBUSY if any files are open.

sfs_open(fileName)
    Open a file, creating it if necessary; returns a small nonnegative “file descriptor” usable only with these routines. Empty names return -EINVAL; failed creates leave no directory entry or allocated block.

sfs_close(fd)
    Close a descriptor. Never fails.

sfs_read(fd, buf, len) / sfs_write(fd, buf, len)
    Read or write up to len bytes at the descriptor’s current position, advancing the position. Writes extend the file (and allocate blocks) as needed.

sfs_remove(name)
    Delete a file. Deleting an open file fails with -EBUSY.

sfs_list(cookie, filename_out, filename_space)
    Iterate over the directory, returning one file name per call.
```

All functions return 0 or a nonnegative result on success and a negative `errno` value on failure. The doc comments in `sfs-api.h` are the authoritative specification, including the exact error codes the driver checks for.

4.3 The functions you must implement

Three functions in `sfs-disk.c` currently return `-ENOSYS`. Implementing them is your first order of business.

```
ssize_t sfs_getpos(int fd)
    Return the current position of descriptor fd, in bytes from the start of the file. If fd is not a valid descriptor, return -EBADF; this is the only possible error. This function is a warm-up to get you reading the descriptor bookkeeping in sfs-disk.c.

ssize_t sfs_seek(int fd, ssize_t delta)
    Move fd’s position by delta bytes and return the result. Clamp it to [0, file size], unlike lseek(2). The descriptor caches both byte position and current block, so crossing block boundaries requires walking the chain and keeping both fields consistent. Given a prior sfs_getpos result loc, sfs_seek(fd, loc - sfs_getpos(fd)) must return the descriptor to the earlier position.

int sfs_rename(const char *old_name, const char *new_name)
    Rename a file. If new_name already exists, atomically delete and replace it: there must be no window in which a concurrent thread can observe new_name not existing. For the single-threaded version this is straightforward; once you add locking, keeping this promise is part of the design work. Renaming an existing file to itself is a no-op; empty names return -EINVAL.
```

5 Making the File System Concurrent

Once the three functions above work, the second half of the lab is to make the file system safe and fast under concurrent use. We suggest approaching this in three gradations.

First, make the system thread-safe in a very coarse way, ensuring correctness only. This can be done by locking the entire file system around every API call, exactly as the calibration reference `sfs-baseline-ref.c` does. Expect a performance ratio near 1.0× and a performance score of 3/10 in the optional benchmark (Section 6.3). A correct global-lock implementation completes the basic exercise; optimization is a follow-up.

Second, think about readers and writers. When reading, no shared data changes except the reading descriptor's own position, so multiple readers of the same file can proceed together, whereas writing wants exclusivity. Be careful about where a reader/writer lock actually helps: because `sfs_read` advances the descriptor's position, a single global rwlock buys almost nothing — nearly every operation still needs the write side. Reader/writer locks pay off for structures that are genuinely read-mostly, such as the directory during name lookups, usually in combination with the per-file locks below.

Third, exploit separate files. Operations on different files touch disjoint data (apart from the shared allocator and tables), so reading or writing one file should not block work on another. A lock per directory slot (or per open file), plus short critical sections on the genuinely shared structures — the free list, the directory, and the open-file tables — is the level of granularity the performance score is calibrated around.

Whatever design you choose, fix a global lock *order* and stick to it everywhere; two threads that take the same locks in opposite orders will deadlock. The driver's 30-second per-trace timeout will catch a deadlock, but a design that cannot deadlock is far more pleasant to debug.

Optional challenges: Zero-block empty files, expandable directories, and Unix unlink are unscored.

6 Evaluation

The basic exercise requires 12 correctness points. The optional benchmark adds 10 performance points, with 4 style points for self-review:

Category	Points	What it tests
A: feature traces (A00–A04)	5	format/mount, open/close/read/write, getpos, seek, rename
B: sequential traces (B00–B03)	4	remove/list and slot reuse, seek/overwrite/append integrity, ENOSPC/EMFILE rollback, name/size edges, rename/list
C: concurrent traces (C00–C02)	3	concurrent writes to separate files, concurrent reads of a shared file, mixed read/write, an open/close storm, and directory churn (create/rename/remove) under a live listing
Performance	10	concurrent throughput against a calibrated baseline
Style	4	manual self-review (not auto-graded)

6.1 Correctness (12 points)

Each trace is worth one point; the driver prints a PASS/FAIL line per trace followed by the details of any failed checks. After every trace unmounts its disk image, the driver also runs `sfs-fsck` on the image: a trace fails if the visible API results look right but the on-disk structure is corrupt (orphaned blocks, inconsistent links, sizes that disagree with the allocation). The first checker diagnostic is included in the failure output.

Category C makes genuinely concurrent calls, including on systems without ThreadSanitizer. The test driver does not supply locks for your code. Failures on the unlocked starter may vary between runs. Start with one mutex around the API, then refine only after correctness passes. The listing trace always finishes each directory walk before starting another, so holding a directory lock across an entire walk is supported.

Every trace — A and B included — runs in a forked child with a 30-second wall-clock budget. A deadlock, crash, or assertion failure in your code therefore fails that one trace (reported as TIMEOUT or CRASH) instead of taking down the driver.

6.2 Race detection

After all twelve correctness traces pass, the driver attempts a ThreadSanitizer build and reruns Category C under three schedule-fuzz seeds. Data races and functional failures fail the concurrency check. The sanitizer has a 120-second budget per sweep, covering the three 30-second traces plus startup overhead. Failure diagnostics remain in `tsan_output.log`.

If the runtime cannot start, it is reported as unavailable without deducting points. Normal concurrent tests still run, but passing tests cannot prove that no race exists. Use a TSan-capable environment for additional checking.

6.3 Performance (10 points)

The optional performance benchmark runs with `make grade` or 8 threads, each working on its own file. A thread performs 16000 *sessions*; one session is one `sfs_open`, then 10 rounds of write/get-pos/seek/read, then one `sfs_close` — 5,376,000 API calls per run in total. Timing starts immediately before a start gate releases the already-created threads, so thread startup is excluded; the benchmark is sampled 5 times and the median ops/sec is scored.

The workload validates its own results as it runs: every return value is checked against the position the descriptor must be at, and every round seeks back and compares the exact bytes it just wrote. If any check fails, the performance score is 0 — an implementation that is fast but wrong on the benchmark path earns nothing.

Raw throughput is machine-dependent, so the score is a ratio against the provided baseline (the hand-out implementation under one global mutex), measured on *your* machine by `make baseline`:

Ratio	Score	Interpretation
≥ 2.50	10/10	scaling that requires fine-grained locking
≥ 1.80	9/10	strong parallelism
≥ 1.40	7/10	meaningful parallelism
≥ 1.20	5/10	measurably better than one global lock
≥ 0.85	3/10	coarse-lock territory (this <i>is</i> the baseline)
< 0.85	0/10	slower than the naive baseline

These thresholds are local comparison targets, not verified official grades. Core count, scheduling, and the machine affect achievable speedup. Use the ratio to compare your own designs; no speedup threshold is required to complete the basic exercise.

If `.perf_baseline` is missing, uses another workload version, or records a different `SFS_DISK_DIR` than the scored run, the driver falls back to absolute throughput *capped at 5/10* and tells you to re-run `make baseline`. Full performance credit always requires a calibrated ratio.

6.4 Style (4 points)

The 4 style points are a self-review: use the included `.clang-format`, remove dead code and debug output, explain *why* instead of what, clean up error paths, and keep locking locally auditable.

7 Testing Your Work

7.1 The test driver

Run `./test-sfs` from the handout directory. `-q` hides FAIL detail; `-v` removes its default cap. Performance never affects the exit status: it is 0 when all twelve correctness traces pass, otherwise 1.

In a non-quiet run, when a trace fails, the driver saves a copy of its surviving disk image(s) under `fail_<trace>_<image>.img` and prints the path, so you can inspect the exact corrupted state instead of re-running under a debugger. The copies live in a private `sfs-test.*` directory, printed by the driver. Remove that directory yourself after inspection, or set `SFS_KEEP_FAILED_DISKS=0` to disable them.

7.2 The consistency checker

`sfs-fsck` verifies the structural integrity of a disk image: mislabeled blocks, invalid directory entries, file sizes that disagree with the number of allocated blocks, inconsistent or circular block chains, blocks on more than one chain, and orphaned blocks on no chain at all.

The driver names its working images `test.img` (Categories A and B), `test_conc_*.img` (Category C), and `test_perf.img` (the benchmark), inside a private `sfs-test.*` directory. It removes working images by the end of a run, preserving failure snapshots. The image you point the checker at is therefore usually a `fail_*.img` snapshot from Section 7.1 — or any image your own experiments create with `sfs_format`:

```
linux> ./sfs-fsck sfs-test.XXXXXX/fail_B01_test.img
linux> ./sfs-fsck -v sfs-test.XXXXXX/fail_B01_test.img
linux> ./sfs-fsck --dump sfs-test.XXXXXX/fail_B01_test.img
```

`--dump` prints what the image actually contains — superblock summary, free list, every directory entry with its block chain, and a block-type tally — before judging it. The dump is deliberately defensive: chains are walked with a step cap and out-of-range block numbers are annotated rather than crashing, so it stays useful on precisely the corrupted images you will want to point it at.

7.3 Race detection by hand

The driver runs TSan for you, but reruns with full stack traces are often more useful when hunting a specific race:

```
linux> gcc -std=c11 -g -fsanitize=thread -pthread -D_GNU_SOURCE=1 \
        -I. -o test-sfs-tsan test-sfs.c sfs-disk.c sfs-support.c
linux> ./test-sfs-tsan --tsan-only --sched-fuzz=2
```

Replace 2 with whatever seed the driver reported. The same seed requests similar perturbations; OS scheduling can still change the execution, so rerunning need not reproduce a failure. The normal correctness pass never fuzzes; only the TSan sweep and explicit `--sched-fuzz` invocations do.

7.4 Calibrating and reading the performance score

Calibrate once per machine with `make baseline`, then run `./test-sfs --benchmark`; `make grade` does both. The calibration prints a spread line; if min and max differ by more than about 10%, the machine was not quiet — close the browser and re-run. Use `make baseline BASELINE_RUNS=11` for a tighter (odd) sample count on noisy machines.

7.5 Notes for Docker and WSL users

Bind-mounted disk I/O can distort the benchmark. Run `export SFS_DISK_DIR=/tmp` before both `make baseline` and `./test-sfs --benchmark`. The driver records the setting in `.perf_baseline`; a mismatch disables ratio scoring and caps performance at 5/10.

8 Hints

- Implement `sfs_getpos` first — it is nearly free once you find the right field — then `sfs_seek`, then `sfs_rename`. Re-run `./test-sfs` after each one.

- Read the block-walking loops in `sfs_read` and `sfs_write` before writing `sfs_seek`; your seek must maintain the same descriptor invariants those loops rely on (position and cached current block agreeing with each other).
- Start the concurrency work coarse, verify correctness, then refine. A correct slow file system outscores a fast corrupt one by a mile: the performance benchmark does not even run until correctness is 12/12.
- Run `sfs-fsck` after concurrent tests. Races often corrupt the block chains in ways that reads happen not to notice; the checker notices.
- Be careful with the operations that touch global state: the free list, the directory, and the open-file tables. Each can have its own small lock, but fix a lock order.

Appendix: Frequently Asked Questions

“No space left on device”

Each test sizes its disk image appropriately for correct execution. If your implementation leaks blocks or double-allocates under a race, an operation can fail with `-ENOSPC` even though a correct implementation would have room. Run `sfs-fsck` on the failing image; leaked blocks show up as “not on any block list.”

How do I debug disk corruption?

`./sfs-fsck -v` narrates the whole check, and `--dump` shows the structure it found. The `fail_*.img` snapshots the driver keeps for failing traces are exactly the images to inspect.

What if `sfs_seek` goes past the end of the file?

Clamp to the file size; if it would go negative, clamp to 0. This is specified in `sfs-api.h` and is deliberately different from `lseek(2)`.